



Pilot Kafka Service

Manuel Martín Márquez



Kafka

- Kafka is a distributed streaming platform
 - High Scalable (partition)
 - Fault Tolerant (replication)
 - Allow high level of parallelism and decoupling between data producers and data consumers
- De facto standard for near real-time store, access and process data streams
- Critical component of most of the Big Data Platform and therefore of Hadoop ecosystem

Kafka Basic Concepts

Broker: Kafka node on the cluster

Topics: Stream of records category

- Multiple writers and readers
- Partitioned
- Replicated

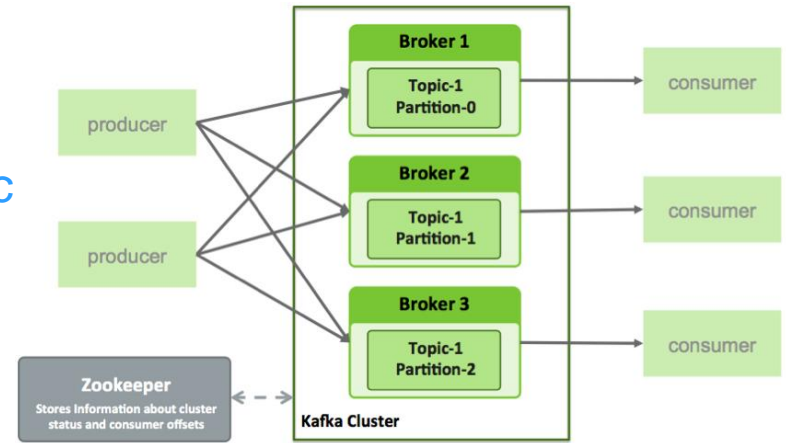
Consumer: pulls messages off of a Kafka topic

Producer: push messages into a Kafka topic

Data Retention:

- Based on time or size

Zookeeper: Stores Kafka Metadata

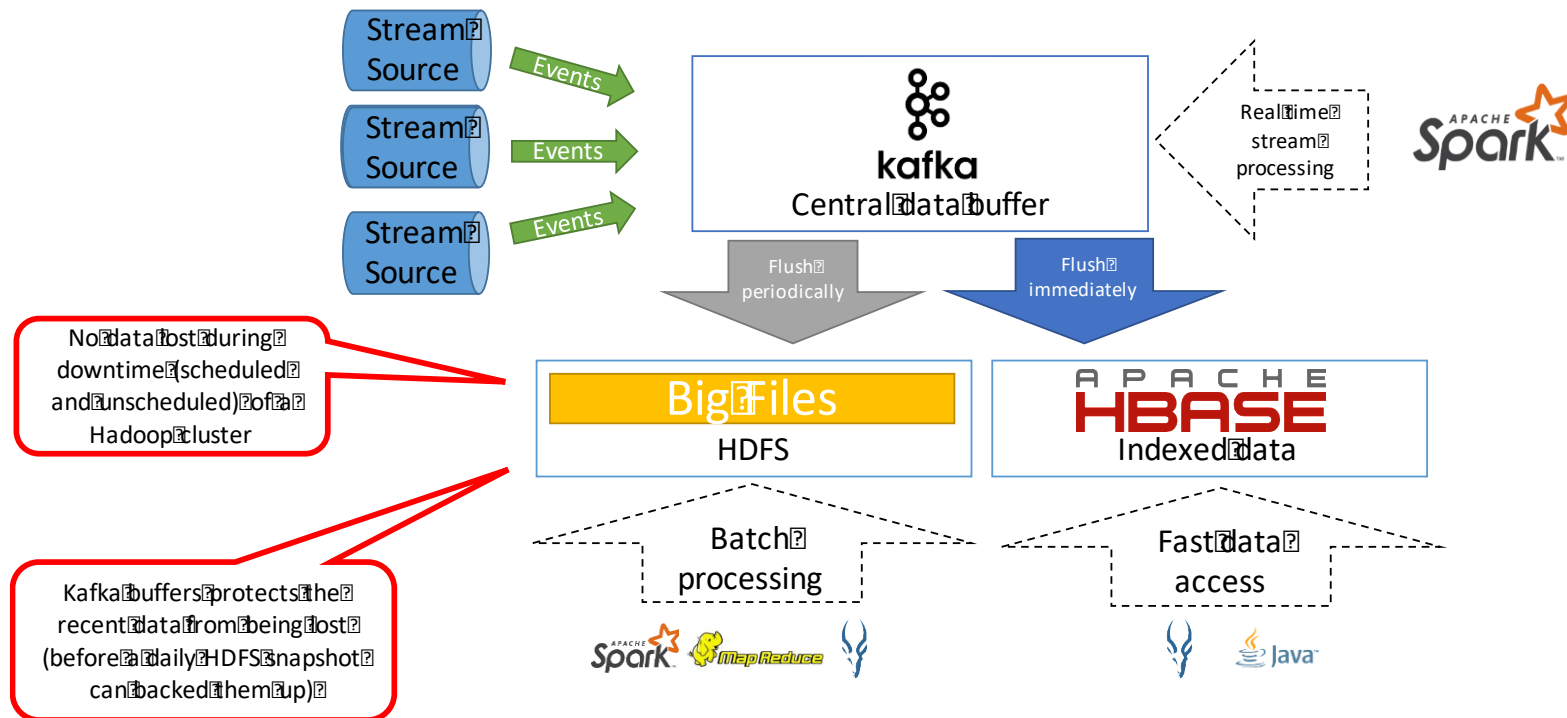


Source: Hortonworks

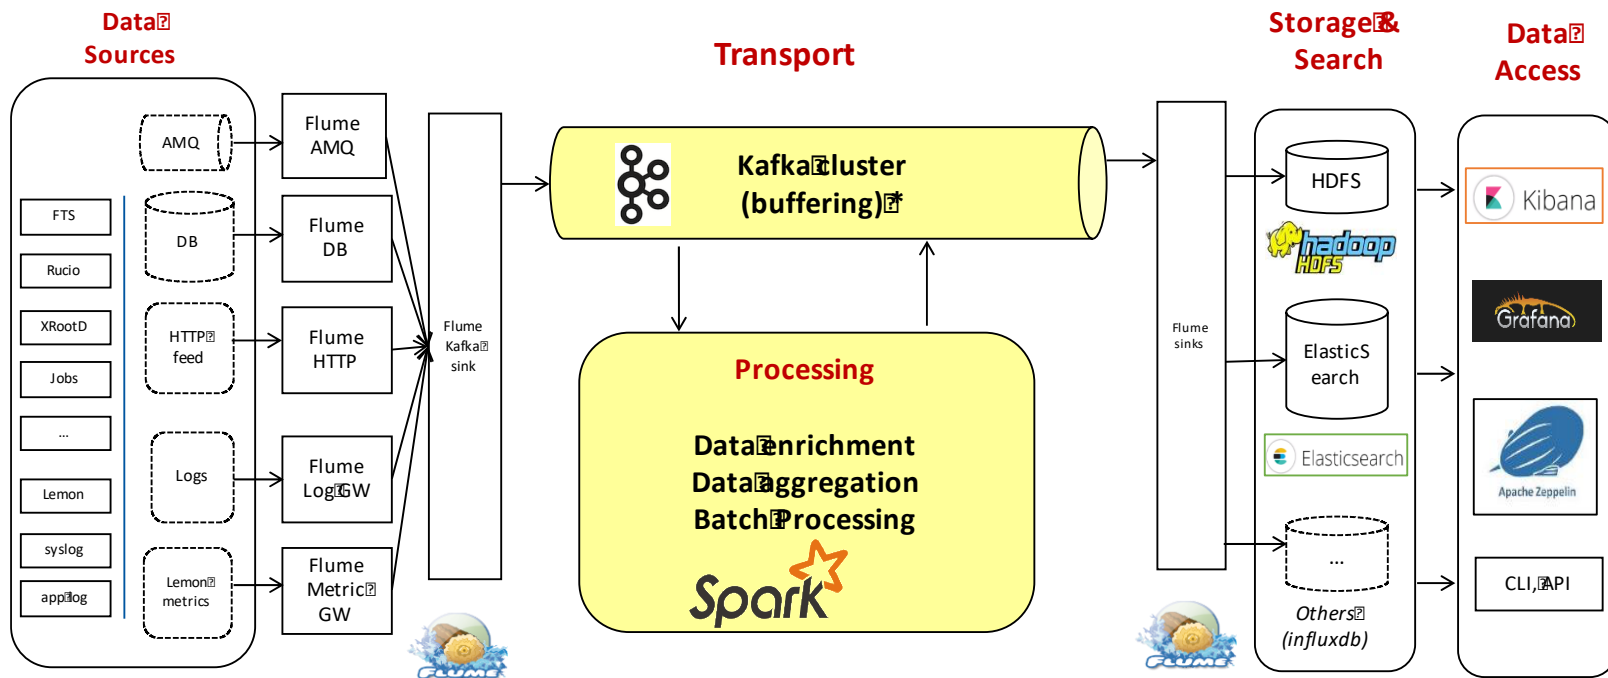
Kafka entry points

- Custom implementation of producer and consumer using Kafka client API
 - Java, Scala, C++, Python
- Kafka Connectors
 - LogFile, HDFS, JDBC, ElasticSearch...
- Logstash
 - Source and sink
- Apache Flume out-of-the-box can use Kafka as
 - Source, Channel, Sink
- Other ingestion or processing tools support Kafka
 - Apache Spark, LinkedIn Gobblin, Apache Storm...

Kafka for Data Integration and Processing



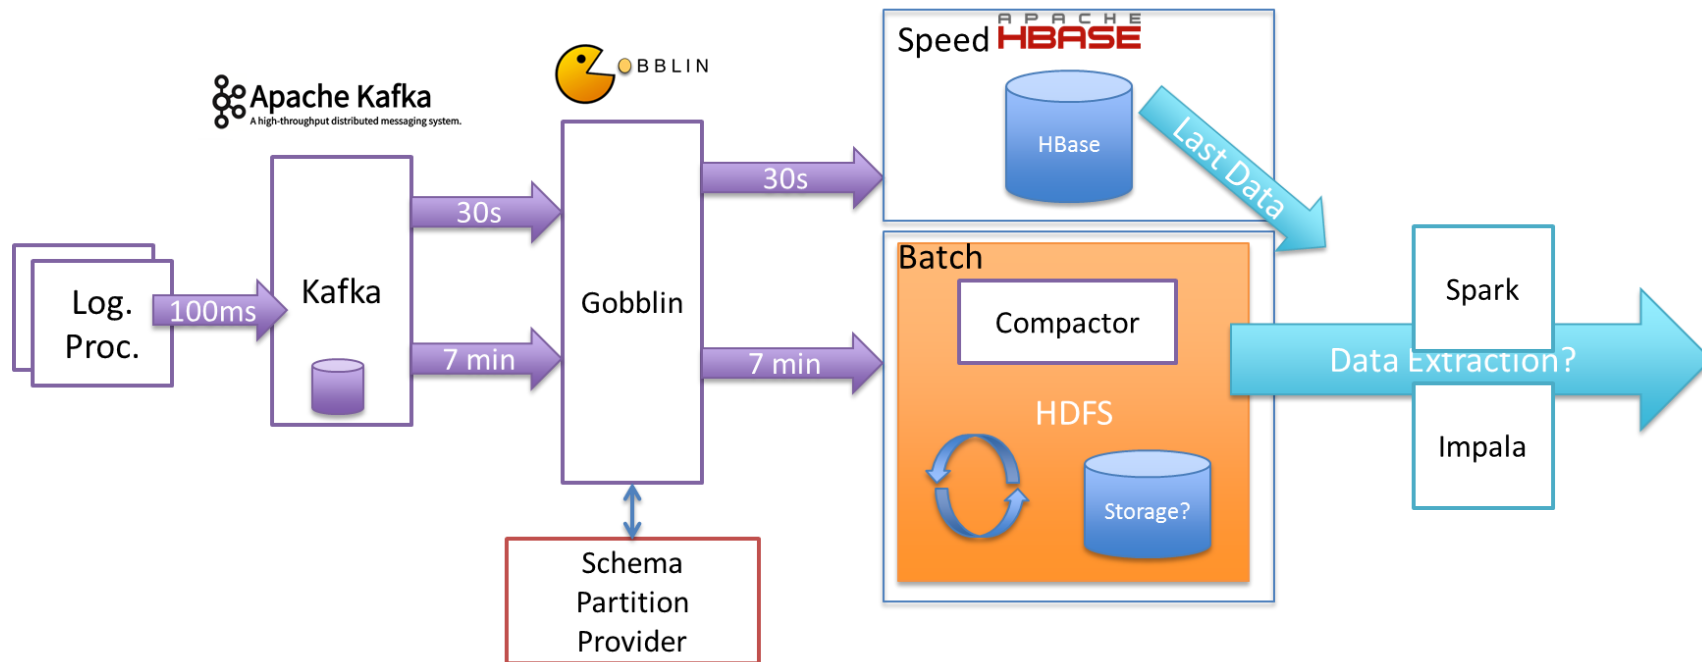
Kafka at CERN – it monitoring



Kafka at CERN – it monitoring (Requirements)

- **Throughput and retention policy**
 - Currently 200 GB/day (forecast 500 GB/day)
 - Retention Policy 12h in qa and 24 hours in prod (largest retention policy to cover potential problems over weekends)
 - ~ 4000 messages, up to 10k peaks
 - ~50 topics
- **Security (Kerberos)**
 - Flume can be potentially upgrade to 1.7 early in 2017 (work in progress already).
- **Administration Capabilities**
 - Administrative operations
 - Topic configuration, rebalancing, user management, start/stop cluster
 - Possibility to increase retention policy, replication factor

Kafka at CERN – CALS



Kafka at CERN – CALS (Requirements)

- **Throughput and retention policy**
 - Currently 30 GB/hour only including the logging processes
 - Plan to incrementally include all the systems with potentially mean several TBs
 - Compression with Snappy will be evaluated to determined performance
 - Retention policy 24 hours, which is the time they need to buffer data and compact it to send it to Hadoop
- **Security (Kerberos)**
- **Infrastructure**
 - Openstack under several conditions:
 - TN need to be supported for several reasons
 - High availability of the service CALS on top of private cloud (No CALS no BEAM in the LHC)
- **Administration Capabilities**
 - Administrative operations
 - Topic configuration, rebalancing, user management, start/stop cluster
 - Possibility to increase retention policy, replication factor

Kafka at CERN

- Security Team
 - Already using Kafka for pattern matching
 - Data integration
- LHC Postmortem
 - Potentially ingested by CALS
- Industrial Control Systems
 - WinCCOA Data

Pilot Kafka Service

- Scope
 - Study the current Kafka use case together with the different teams involved
 - Collect requirements
 - Understand feasibility and added value of Kafka as a central service

Pilot Kafka Service

- Collect requirements (5 Major Use Cases):
 - **CALS, IT-Monitoring**, Security Team, Industrial Control, Post-mortem
 - Throughput, Retention Policy, Security, Infrastructure, Administration Capabilities
 - Agreement to test the service from the first phase
 - Ensure the service cope with their requirements
 - More details: <https://twiki.cern.ch/twiki/bin/viewauth/DB/CERNOnly/KafkaService>

Pilot Kafka Service — Current Development

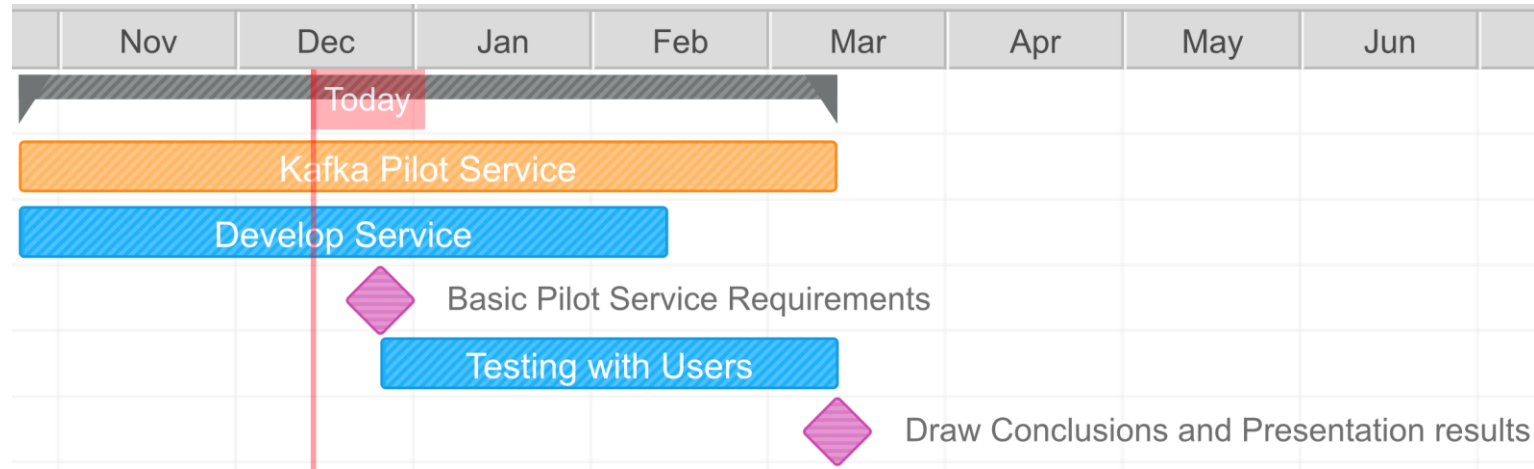
- Pilot Implementation - rapid iteration which will help to understand service and use case.
- On-demand Kafka service approach
 - Self-Service Cluster creation, management and expansion
 - Allow users to perform administrative tasks that are traditionally carried out by administrators
 - Facilitating operating system and engine updates (Kafka, Zookeeper)
 - Transparently integrate all the needed services (Security, Storage, Procurement, etc)
 - Support for service continuity in case of hardware failure

Pilot Kafka Service — Current Development

- Configuration and Management REST API
- Security enabled - Kerberos on Kafka and Zookeeper (SSL optional)
- Monitoring Capabilities
- OpenStack on GPN
- Network storage
- Dedicated Kafka and Zookeeper per user

Towards Kafka Production Service

- Service evaluation phase and time line



Towards Kafka Production Service

- Consolidation to Production
 - Web Interface to manage clusters (Self-service)
 - Evolution of the configuration management API
 - Functionalities toward the self-service platform
 - Integration with Openstack
 - Full monitoring beyond JMX metrics
 - Kafka-Mirroring (High Availability)
 - Deploy service in TN (due to service design that is transparent for us)
 - Kafka as close as possible to consumers and producers