

# Introduction to MongoDB



# Challenges with Relational Databases



Adding features to relational databases becomes challenging due to complexity.



Users face difficulties understanding and navigating applications with multiple tables.



As features grow, maintaining efficiency in a bloated, inefficient table is problematic.



The relational approach poses challenges for developers and application users.



Pulling data from multiple places in a relational database results in inefficiency, impacting both speed and user experience.

# Relational Database

Traditional relational databases resemble Microsoft Excel on an enhanced scale.



Rows and columns organize data into tables, forming a rigid structure.



This Excel-like structure creates challenges for evolving data requirements.

# Introduction to MongoDB



MongoDB's document model stores data in flexible BSON (Binary JSON) documents.



This shift is fundamental, diverging from traditional relational databases.



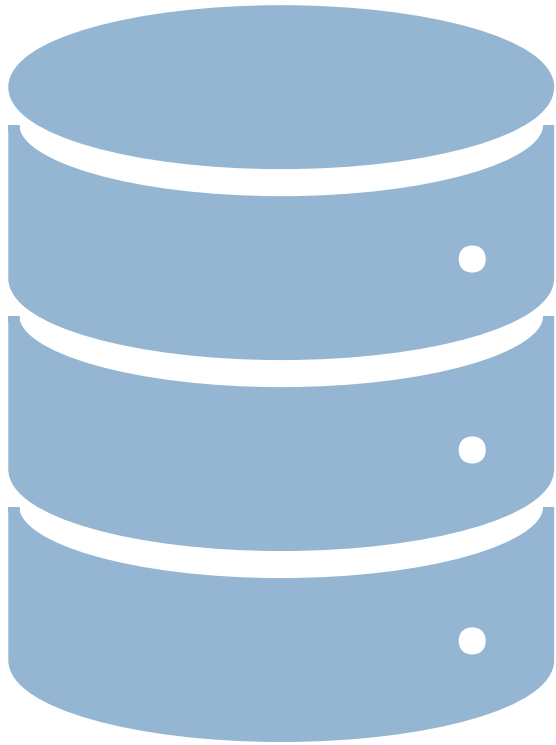
The document-oriented structure is a key design choice, defining MongoDB.



Developers appreciate MongoDB's approach, making it a popular choice.



Let's explore how this model differs in practice.



# Introduction to MongoDB

- MongoDB is a NoSQL document database.
- Data is organized as documents, further grouped into collections in MongoDB.
- Scalable High-Performance Open-source, Document-orientated database.
- Built for Speed
- MongoDB is written in C++, JavaScript, and C language

# Introduction to MongoDB



Rich Document based queries for Easy readability.



Full Index Support for High Performance.



Replication and Failover for High Availability.



Auto Sharding for Easy Scalability.



Map / Reduce for Aggregation.



# Document Model Overview

- MongoDB's document-oriented structure is the cornerstone of its design.
- Data is stored in BSON documents, offering flexibility for dynamic changes.
- BSON allows efficient handling of diverse data types in a natural format.
- The document model aligns with JSON-like structures, enhancing readability.
- MongoDB's approach empowers developers with a user-friendly data model.

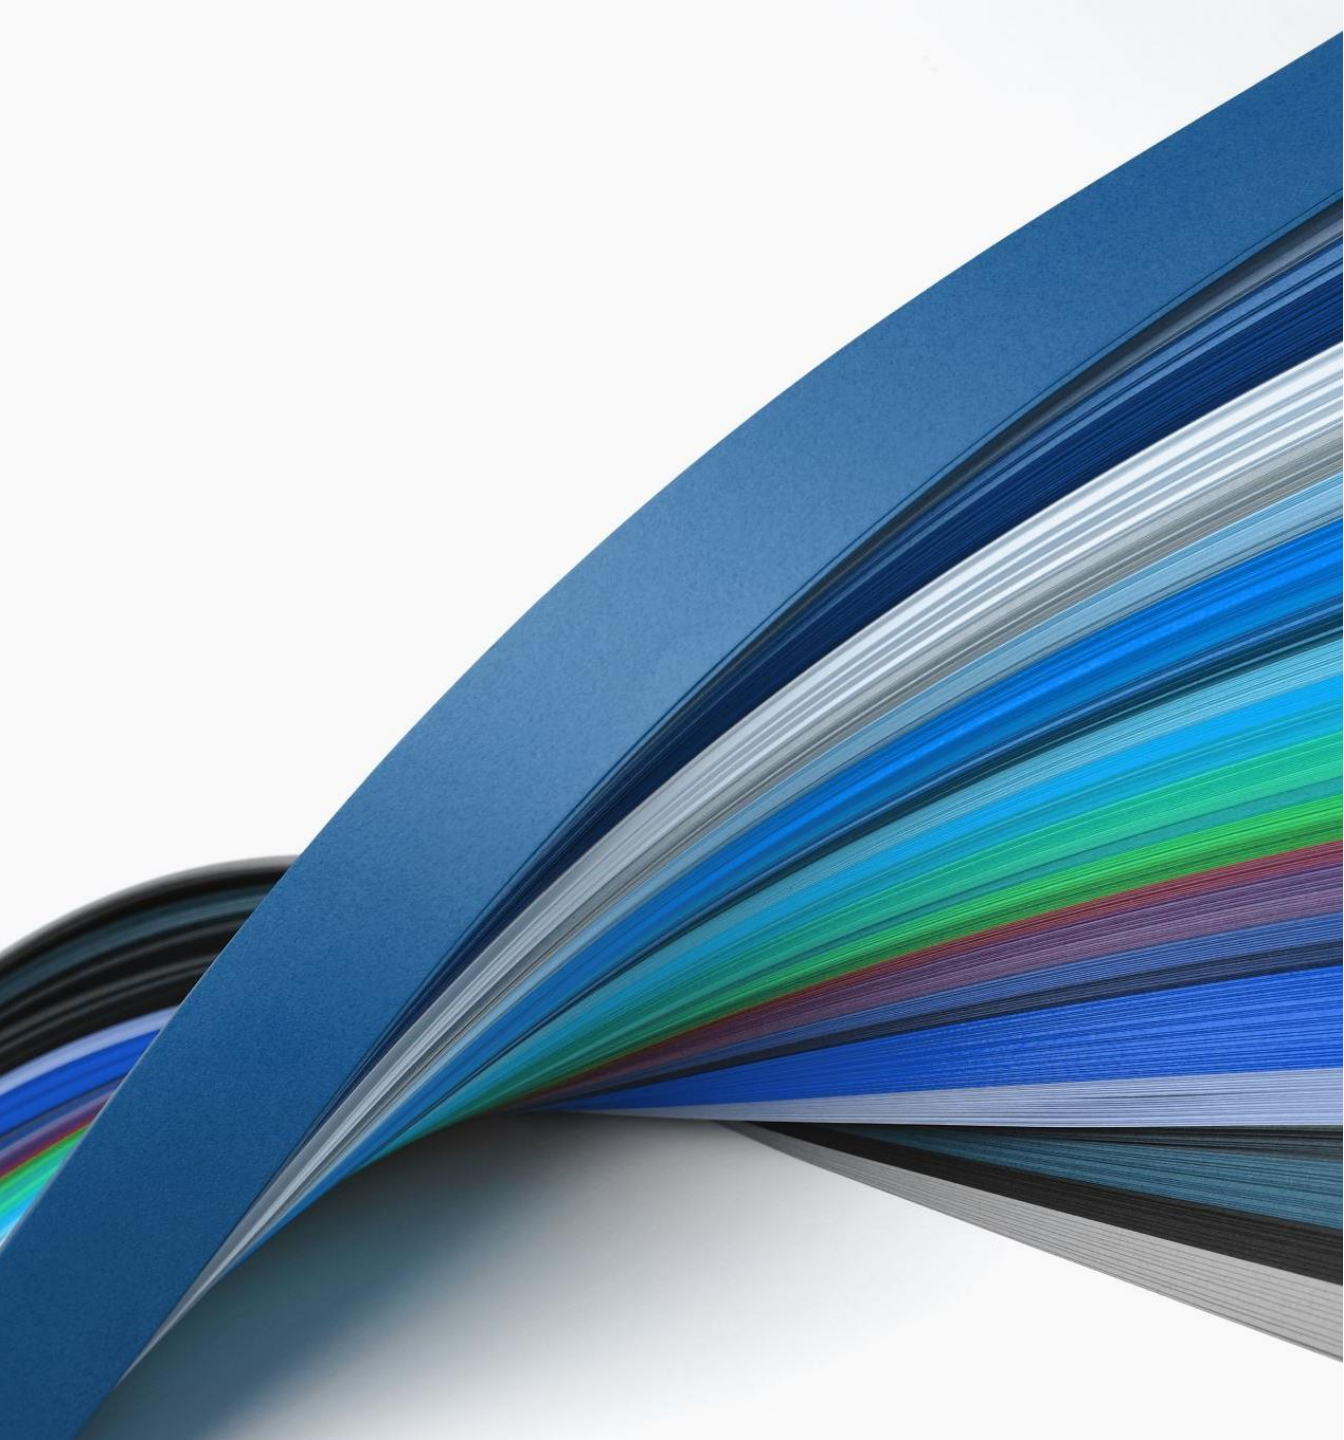




# Advantages of Document Storage

- MongoDB's document storage mimics the structure of physical medical records.
- Data is organized in a way that is both efficient for computers and readable for humans.
- The document model allows for the inclusion of varying data attributes seamlessly.
- Developers appreciate the power of adapting data structures as needed.
- MongoDB's approach facilitates efficient storage and retrieval of diverse data.





# Flexibility for Developers

- MongoDB empowers developers by accommodating the needs of their applications.
- Developers can store data in a natural way without rigid schema constraints.
- This flexibility ensures that MongoDB adapts to the evolving needs of applications.
- MongoDB's document model simplifies the development process.
- Adding or modifying data becomes a straightforward task for developers.



# MongoDB's Distributed System

## Coordination of Multiple Servers

- MongoDB excels in fault tolerance by keeping redundant copies across servers.
- A single server failure does not disrupt the overall application.
- The distributed nature of MongoDB allows for horizontal scaling with ease.
- Data coordination across multiple servers ensures high availability.
- MongoDB's distributed system architecture is a key advantage over traditional databases.



# Fault Tolerance in MongoDB

- MongoDB's native fault tolerance is achieved through redundant data copies.
- In case of a server failure, the system automatically switches to a functioning server.
- This fault tolerance feature enhances the reliability of MongoDB in real-world scenarios.
- Applications remain resilient, even in the face of unexpected server issues.
- MongoDB's fault tolerance is a critical aspect for robust and uninterrupted operations.



# Horizontal Scaling with Sharding

## Sharding for Scalability

- MongoDB's sharding capability facilitates horizontal scaling for data growth.
- Adding more servers is a seamless solution for scaling MongoDB horizontally.
- Sharding ensures that data is distributed effectively across different instances.
- MongoDB's approach to scalability avoids the need for costly hardware upgrades.
- Horizontal scaling is particularly advantageous for cloud environments.

Horizontal scaling involves adding more machines or nodes to a system, while vertical scaling involves adding more power (CPU, RAM, storage, etc.) to an existing machine.

## **Horizontal scaling Vs Vertical scaling**

# MongoDB for Real-time Access

## Efficient Data Access

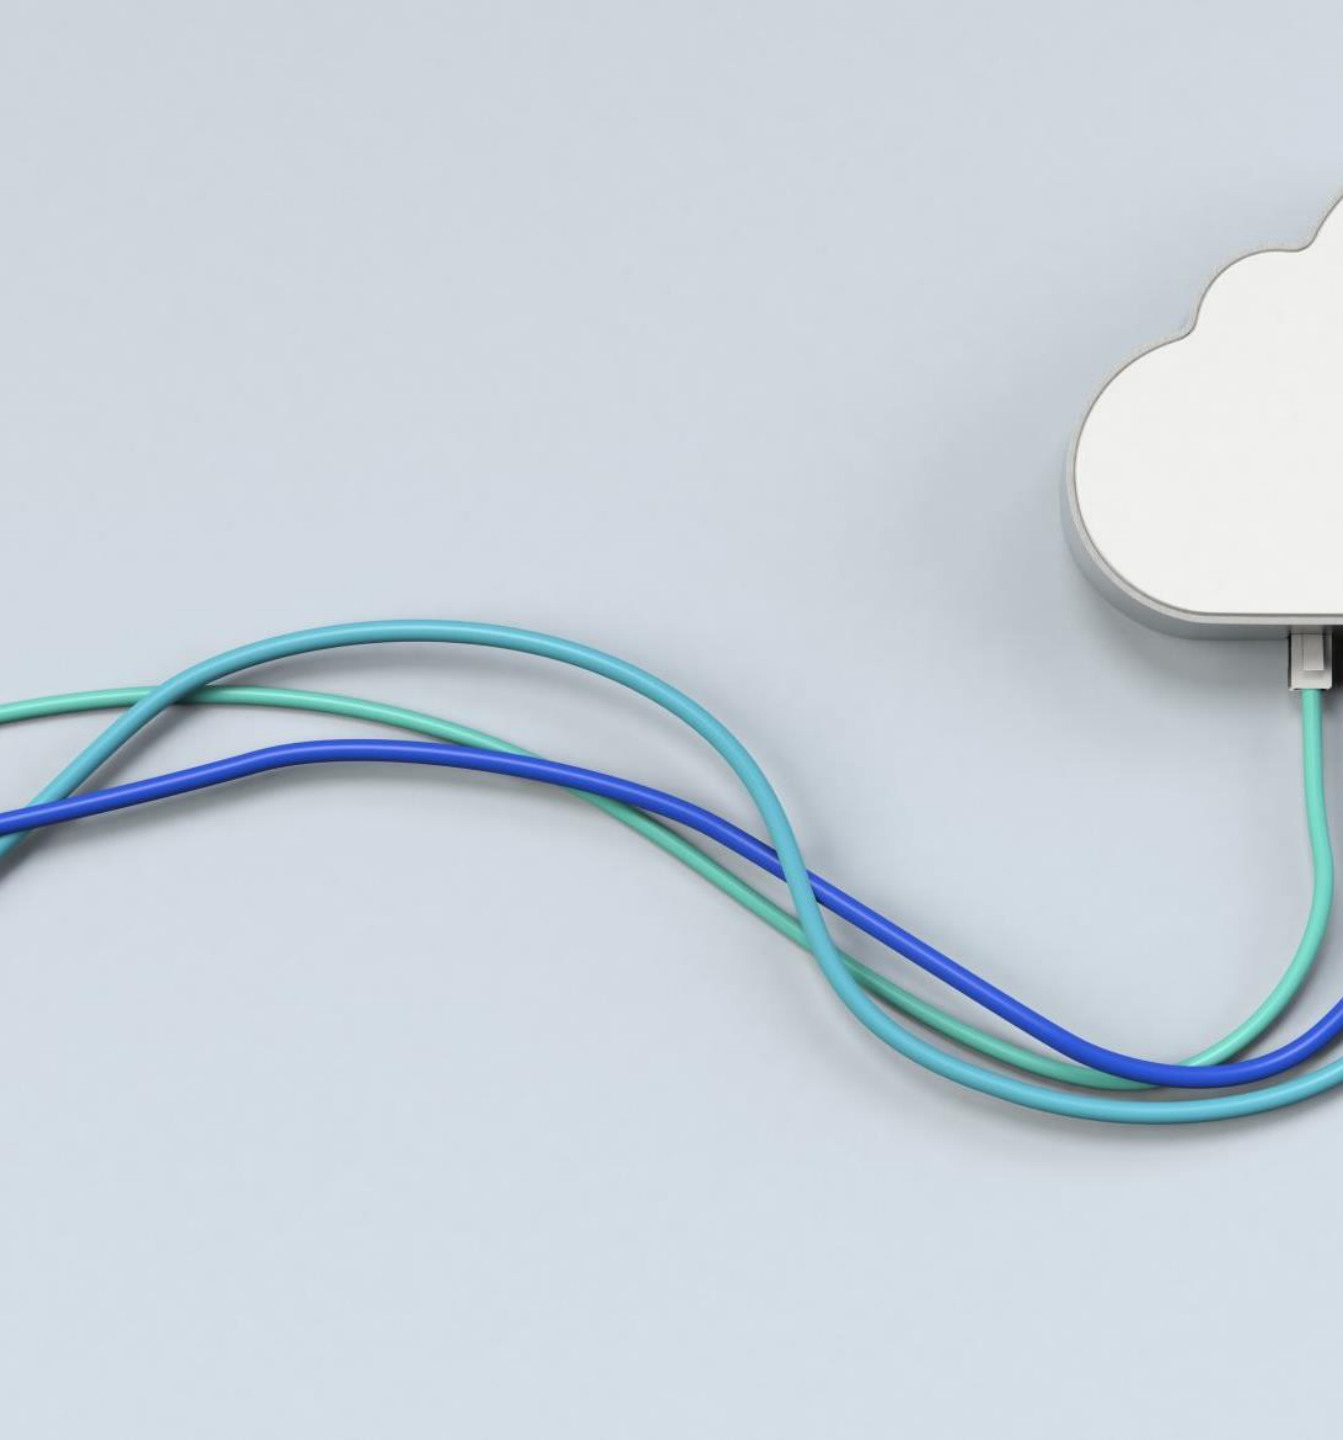
- MongoDB allows data to be moved globally, ensuring fast access for users.
- Real-time access to data is achieved through MongoDB's distributed architecture.
- Data can be strategically placed near users, optimizing access times.
- MongoDB's architecture is well-suited for applications requiring real-time responses.
- Efficient data access is a key strength in MongoDB's design.





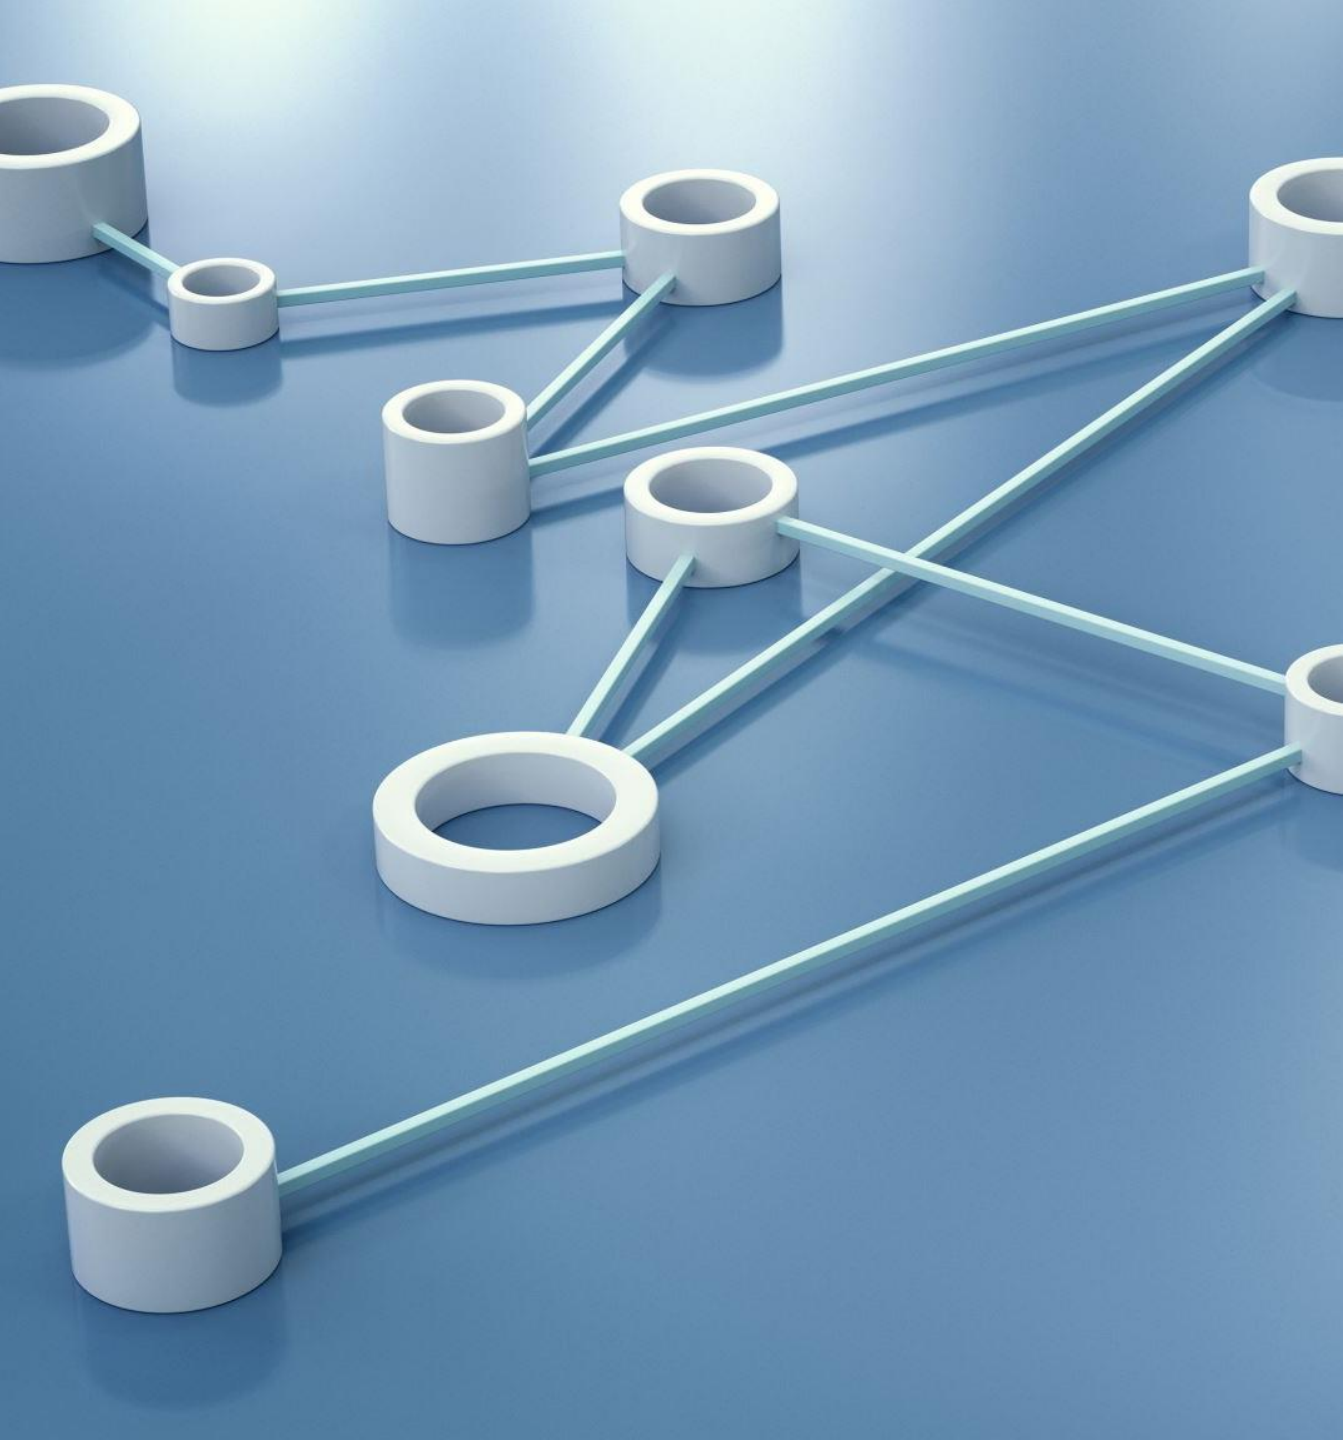
- MongoDB offers a suite of management tools for various aspects of database operations.
- Provisioning tools help set up and configure MongoDB clusters effectively.
- Security features, monitoring, and backup tools are critical for robust database management.
- MongoDB's management tools are tailored to meet the demands of high-performance use cases.
- These tools ensure operational excellence and adherence to strict SLAs.





# MongoDB for the Cloud

- MongoDB is designed to excel in cloud environments, aligning with modern computing trends.
- MongoDB Atlas, the database-as-a-service offering, provides worry-free management.
- Availability on AWS, Azure, and Google Cloud Platform offers flexibility to users.
- MongoDB's cloud-centric design is in line with the preferences of many organizations.
- The move toward cloud-based services is driving MongoDB's widespread adoption.



# MongoDB Atlas

- MongoDB Atlas is a fully elastic database-as-a-service solution.
- It simplifies database management by handling configurations, security, and monitoring.
- Availability on major cloud platforms gives users the freedom to choose providers.
- MongoDB Atlas ensures users can focus on application development without database management worries.
- The platform aligns with the need for scalable and reliable cloud-based databases.



# MongoDB in Real-world Applications

- MongoDB finds applications in renowned companies such as Toyota, SAP, and Adobe.
- Widely deployed in areas like IoT, mobile applications, and real-time analysis.
- MongoDB's versatility makes it a preferred choice for handling large-scale data requirements.
- The adoption by major corporations highlights MongoDB's reliability and effectiveness.
- Real-world success stories emphasize MongoDB's impact on various industries.





# MongoDB Security Features

- Authentication and authorization features.
- Native data encryption for enhanced protection.



## Document

A way to organize  
and store data as a set  
of field-value pairs

```
{  
  <field> : <value>,  
  <field> : <value>,  
  "name"  : "Lakshmi",  
  "title" : "Team Lead",  
  "age"   : 26  
}
```

# MONGODB

# What is MongoDB great for?

---

RDBMS replacement for Web Applications.

---

Semi-structured Content Management.

---

Real-time Analytics & High-Speed Logging.

---

Caching and High Scalability

---

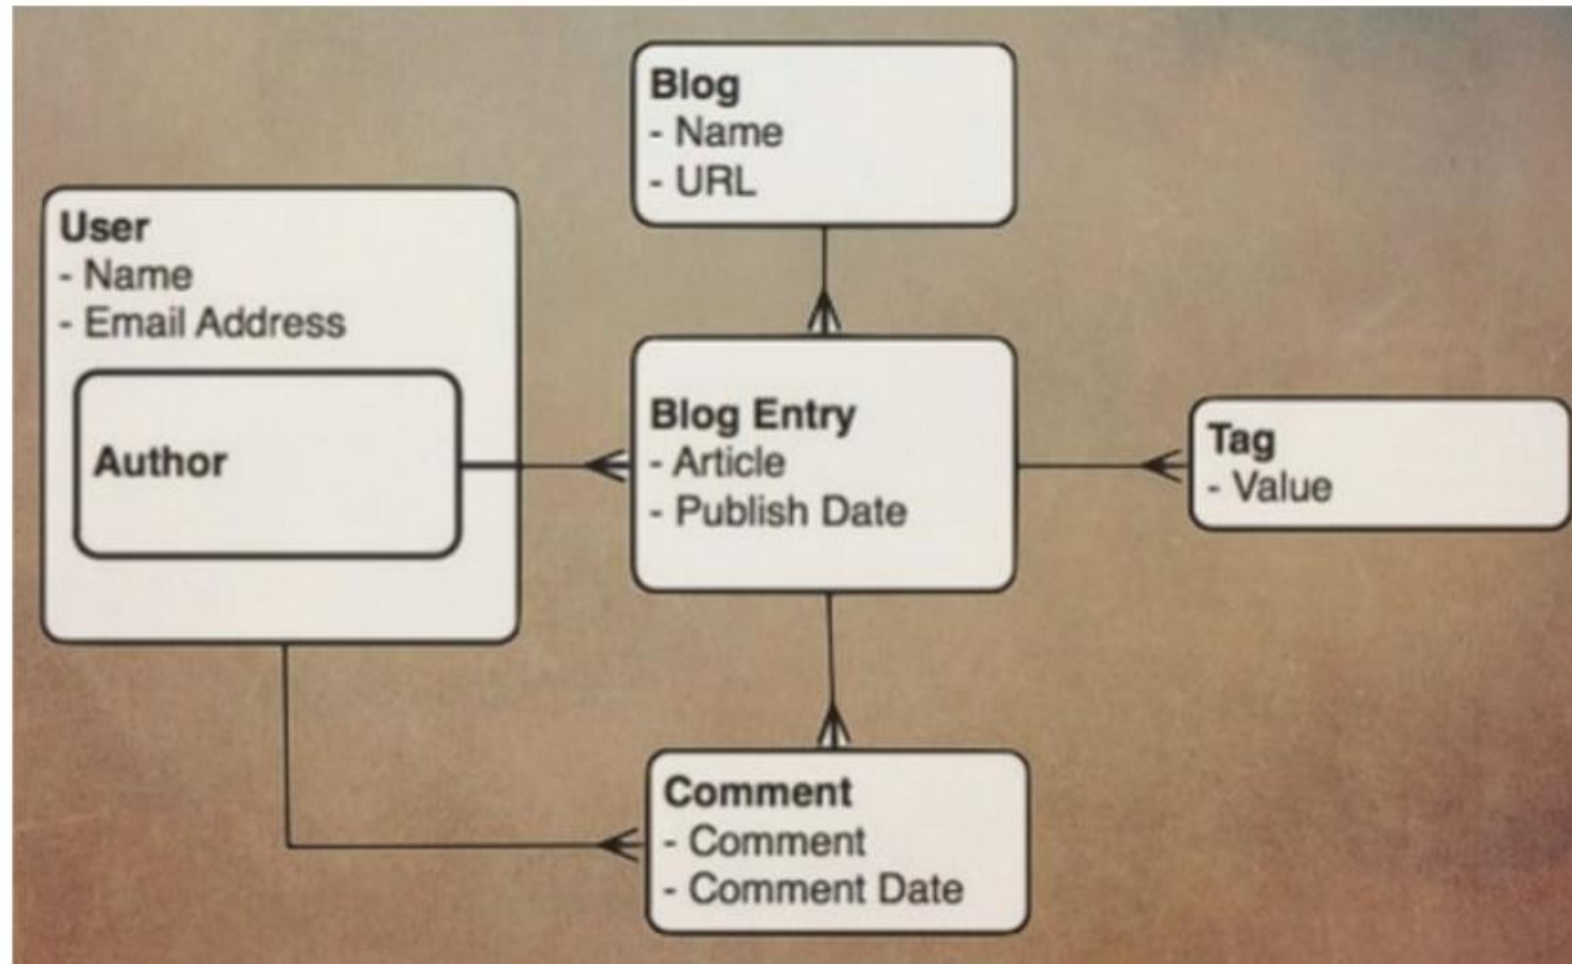
Web 2.0, Media, SAAS, Gaming HealthCare,  
Finance, Telecom, Government

**Let's Dive in !**

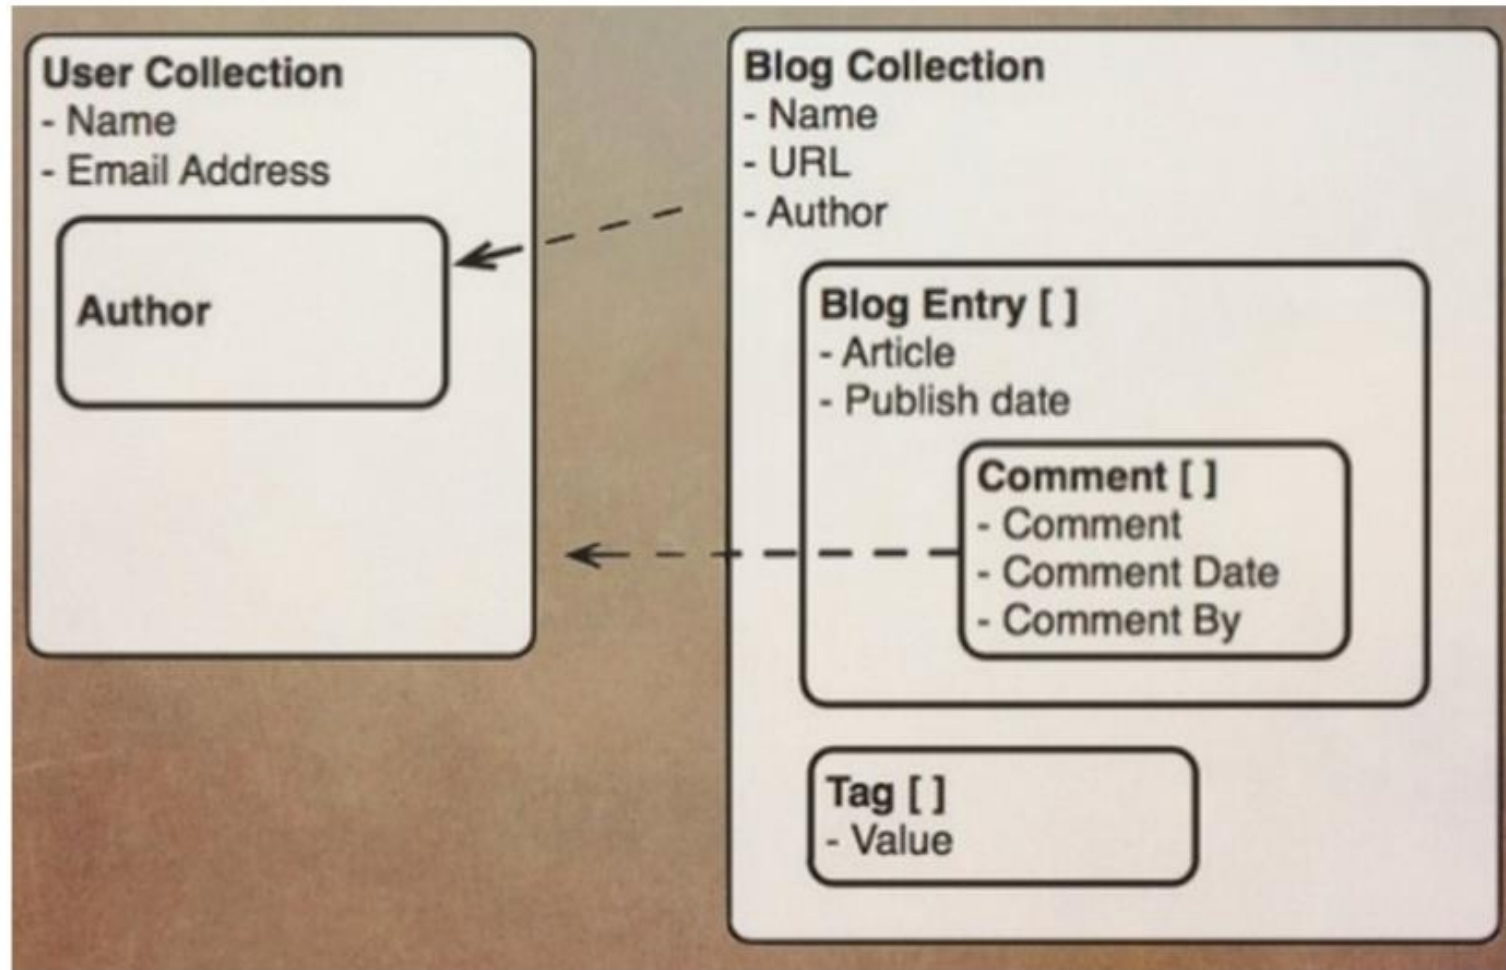




## Blog in relational DB



## Blog post structure in document DB



## Blog post in JSON DB

```
{  _id : ObjectId("4c4ba5c0672c685e5e8aabf3"),
    author : "steve",
    date : "Sat Apr 24 2013 19:47:11",
    text : "About MongoDB...",
    tags : [ "tech", "databases" ],
    comments : [
      {
        author : "Fred",
        date : "Sat Apr 25 2013 20:51:03 GMT-0700",
        text : "Best Post Ever!"
      }
    ]
}
```

**When I say**  
**Database**



**Think**  
**Database**

- Made up of Multiple **Collections**.
- Created **on-the-fly** when referenced for the first time.

When I say  
Collection



Think  
Table

- Schema-less, and contains Documents.
- Indexable by one/more keys.
- Created on-the-fly when referenced for the first time.
- Capped Collections: Fixed size, older records get dropped after reaching the limit.

## Collection

An organized store of documents in MongoDB, usually with common fields between documents

```
{  
  "name" : "Lakshmi".  
  {  
    "name" : "Pavi",  
    "title" : "Engineer",  
    "age" : 21  
  }  
}
```

Query	Description
<code>db.collection.find({})</code>	Retrieve all documents in the collection.
<code>db.collection.find({ field: value })</code>	Retrieve documents matching specified field-value pair.
<code>db.collection.insertOne({})</code>	Insert a single document into the collection.
<code>db.collection.insertMany([{}], {})</code>	Insert multiple documents into the collection.
<code>db.collection.updateOne({ filter }, { \$set: { field: value } })</code>	Update a single document based on a filter.
<code>db.collection.updateMany({ filter }, { \$set: { field: value } })</code>	Update multiple documents based on a filter.
<code>db.collection.deleteOne({ filter })</code>	Delete a single document based on a filter.
<code>db.collection.deleteMany({ filter })</code>	Delete multiple documents based on a filter.
<code>db.collection.aggregate([pipeline])</code>	Perform aggregation operations on the collection.



```
var p = {  
  '_id': '3432',  
  'author': DBRef('User', 2),  
  'title': 'Introduction to MongoDB',  
  'body': 'MongoDB is an open sources.. ',  
  'timestamp': Date('01-04-12'),  
  'tags': ['MongoDB', 'NoSQL'],  
  'comments': [{ 'author': DBRef('User', 4),  
                  'date': Date('02-04-12'),  
                  'text': 'Did you see.. ',  
                  'upvotes': 7, ... }  
]  
}  
> db.posts.save(p);
```



## Understanding the Document Model.

Create Index on any field in the document

// 1 means ascending, -1 means descending

```
> db.posts.ensureIndex({'author': 1});
```

//Index Nested Documents

```
> db.posts.ensureIndex('comments.author': 1);
```

// Index on tags

```
> db.posts.ensureIndex({'tags': 1});
```

// Geo-spatial Index

```
> db.posts.ensureIndex({'author.location': '2d'});
```



## Secondary Indexes

# What about Queries? So Simple

```
// find posts which has 'MongoDB' tag.
```

```
> db.posts.find({tags: 'MongoDB'});
```

```
// find posts by author's comments.
```

```
> db.posts.find({'comments.author':  
DBRef('User',2)}).count();
```

```
// find posts written after 31st March.
```

```
> db.posts.find({'timestamp': {'gte': Date('31-03-12')}});
```

```
// find posts written by authors around [22, 42]
```

```
> db.posts.find({'author.location': {'near':[22, 42]}});
```

```
$gt, $lt, $gte, $lte, $ne, $all, $in, $nin, count, limit, skip, group, etc...
```





### Use Case

Real-time analytics, content management, IOT, mobile apps

### Use Case

Legacy applications or applications that require multi-row transactions

### Data Structure

No schema definition required

### Data Structure

Structured data with clear schema

### Risk

Less risk of attack due to design

### Risk

Risk of SQL injection attacks

### Analysis

MongoDB works great for unstructured data and lends you opportunity for growth.

### Analysis

MySQL is a perfect when your data is structured and you are in need of a traditional relationship database.



craigslist



MetLife

NETFLIX



CITRIX

# MongoDB Applications



Internet of  
Things



Mobile  
applications



Real time  
analysis



Personalization



Catalog  
management

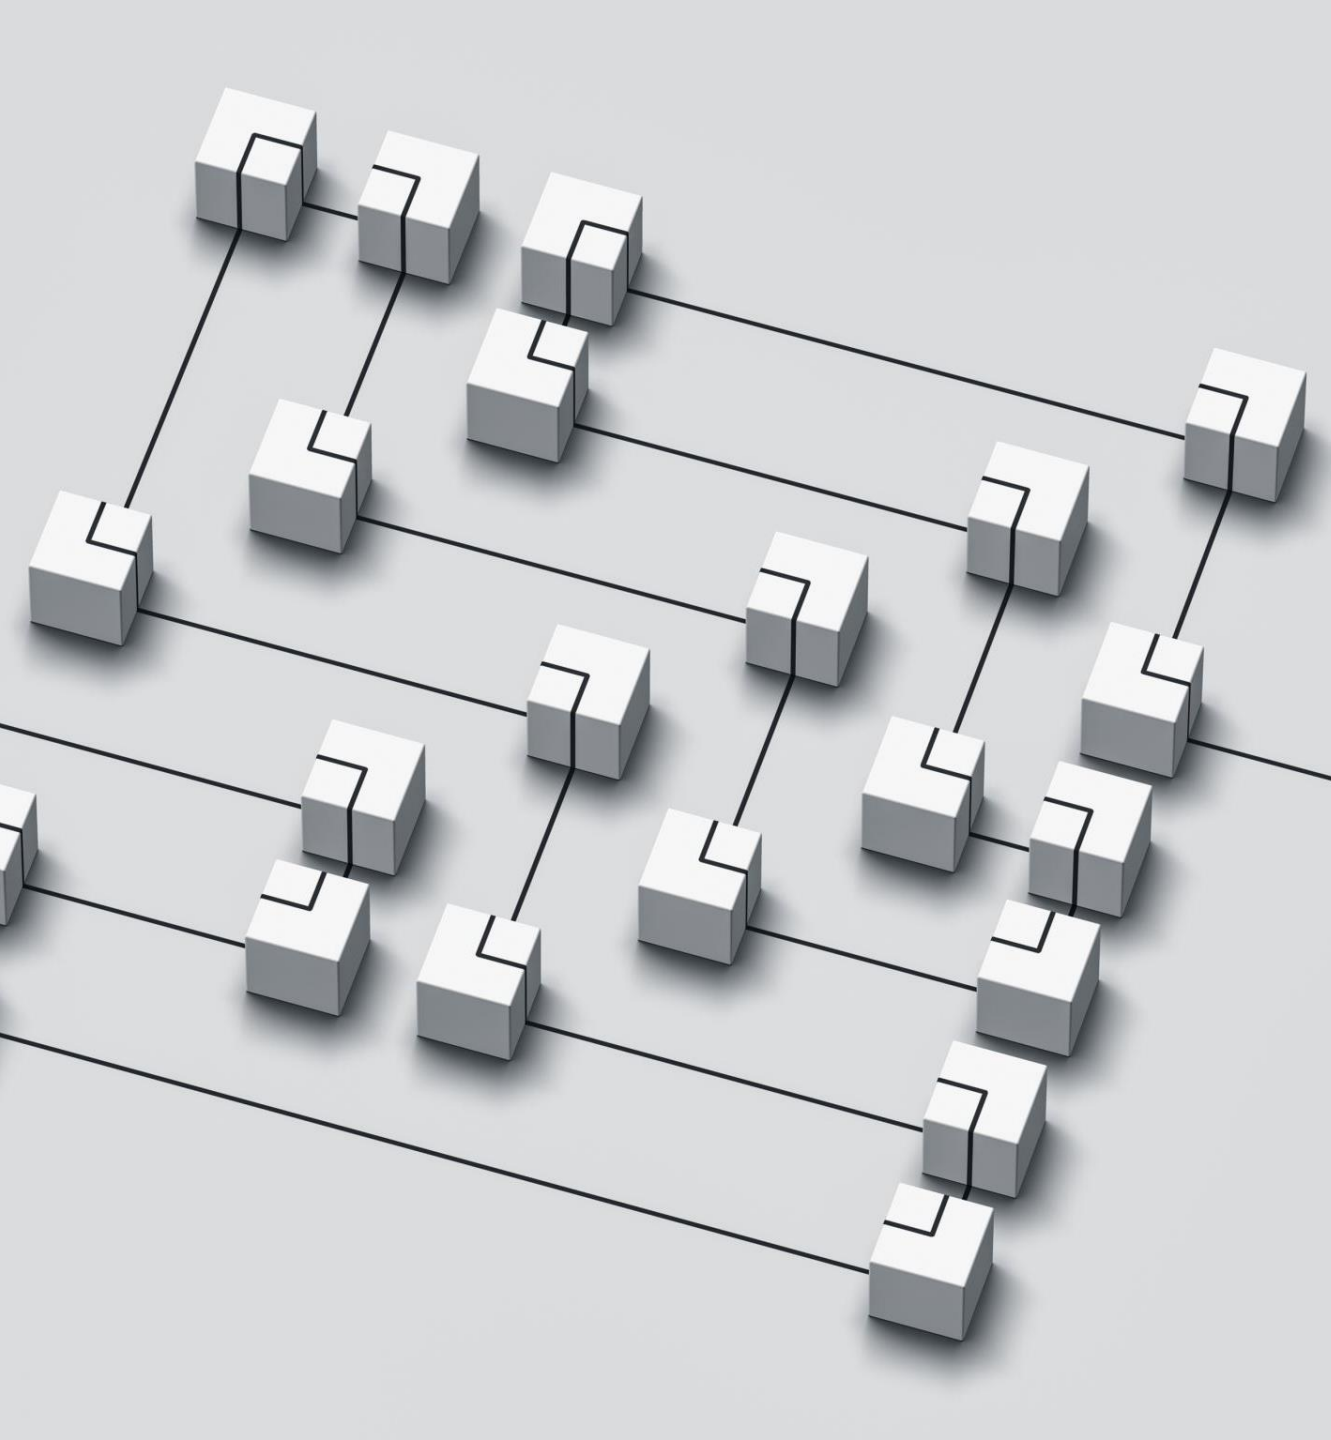


Content  
management

Feature	Relational Database	MongoDB
Data Structure	Tables with rows and columns	Collections with flexible BSON documents
Schema	Fixed schema	Dynamic schema, adaptability over time
Query Language	SQL	MongoDB Query Language (MQL)
Flexibility	Rigid structure	Flexible, accommodates evolving data
Relationships	Complex relationships	Simple and nested structures
Scalability	Vertical scaling	Horizontal scaling with sharding
Complexity	High	Lower, especially for developers
Use Cases	Traditional applications	Diverse applications, big data, IoT
Key Strengths	ACID transactions	High performance, scalability, flexibility
Data Retrieval Efficiency	Joins required	Efficient retrieval, no need for joins
Schema Evolution	Challenging	Seamless, supports dynamic changes
Examples	MySQL, PostgreSQL	MongoDB

# COMPARISON - RELATIONAL DATABASE VS. MONGODB





# MongoDB's Advantages Over Relational Databases

- MongoDB's combination of the document model and distributed systems offers a competitive edge.
- The architecture is ideal for diverse real-world applications with complex data needs.
- MongoDB's approach simplifies development, making it attractive to a wide range of developers.
- The advantages over relational databases are evident in terms of both efficiency and scalability.
- MongoDB's design choices position it as a leading modern database solution.



# MongoDB Use Cases

- MongoDB is applied across diverse technologies, including Internet of Things (IoT).
- Its adaptability makes it suitable for mobile applications and personalized user experiences.
- Real-time analysis benefits from MongoDB's efficient data retrieval capabilities.
- Catalog management and content management are additional areas where MongoDB shines.
- MongoDB's broad range of use cases demonstrates its flexibility and applicability.

# MongoDB Customers

- eBay employs MongoDB for cloud management, search suggestions, and more.
- Shutterfly switched from Oracle to MongoDB for its 6 billion images.
- Aadhar, India's ID project, uses MongoDB for the world's largest biometric database.
- Electronic Arts relies on MongoDB for FIFA Online 3, their online multiplayer game.
- <https://www.mongodb.com/who-uses-mongodb>